


dev-loop-ultraight

One reviewer, one round, no verification phase — a nine-concern sweep in a single pass. Review cost = 0.1x a full run. Implementation is NOT downgraded.




Four consolidated lanes, two rounds, verification only when it is earned. Review cost  $\approx 0.3\times$  a full run. The life loop economises on review, never on the code



Nine gated lanes, one owner per concern, three rounds. Four Opus lanes and five Sonnet. The harvest default for most work — the 1\* against which the other four loops are priced.




The same nine lanes as dev-loop, but each is reviewed twice on opposed assumptions and recorded by a third agent. Roughly 3x the review cost. Implementation is unchanged — ultra buys scrutiny, not code




dev-loop-ultra-opus — adversarial, every subagent on Opus

Structurally identical to dev-loop-ultra. It changes exactly two things: the model on every agent, and the concurrency. Roughly 4x the review cost of a full run, and 2.7x the run as a whole — the most expensive loop available.



Four consolidated lanes plus a Fable Visual lane, two rounds, verification narrower than the full loop. Escalates to dev-loop-unity on a round-1 Critical or a scope break, to the general ladder if eligibility breaks.



Six gated lanes across two rows, a Fable-driven Visual lane, Security dropped since an eligible change has no reachable surface. Two rounds plus a visual confirmation. No heavier Unity loop to escalate to

