



Nine gated lanes, one owner per concern, three rounds. Four Opus lanes and five Sonnet. The honest default for most work — the 1× against which the other four loops are priced.




The same nine lanes as dev-loop, but each is reviewed twice on opposed assumptions and reconciled by a third agent. Roughly 3× the review cost. Implementation is unchanged — ultra buys scrutiny, not code.




dev-loop-ultra-opus — adversarial, every subagent on Opus

Structurally identical to dev-loop-ultra. It changes exactly two things: the model on every agent, and the concurrency. Roughly 4× the review cost of a full run, and 2.7× the run as a whole — the most expensive loop available.



Four consolidated lanes, two rounds, verification only when it is earned. Review cost $\approx 0.3\times$ a full run. The lite loop economises on review, never on the code.



One reviewer, one round, no verification phase — a nine-concern sweep in a single pass. Review cost $\approx 0.1 \times$ a full run. Implementation is NOT downgraded.

